

# Standard I/O in Linux, pipes, redirection, regular expressions, text processing

Nikita Neveditsin, SMU, 2019

[nikita.neveditsin@smu.ca](mailto:nikita.neveditsin@smu.ca)

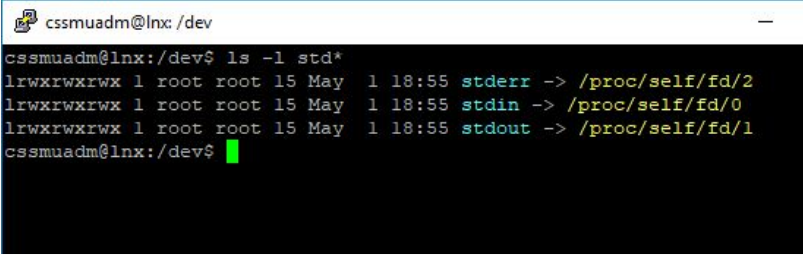


# Standard Input and Output

In Linux all processes are given 3 streams for **input** and **output** of data:

- Standard Input (STDIN) - file descriptor 0
- Standard Output (STDOUT) - file descriptor 1
- Standard Error (STDERR) - file descriptor 2

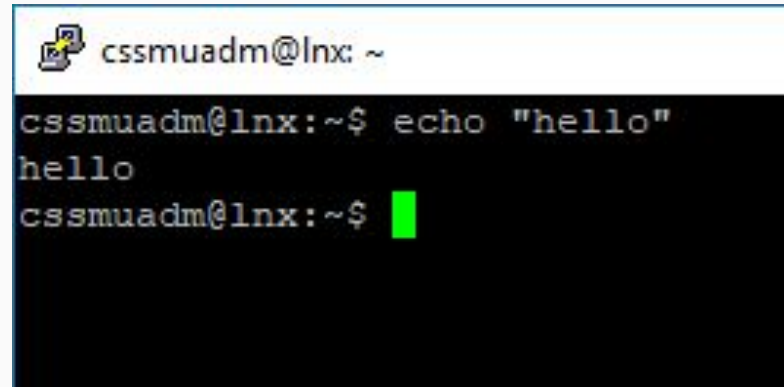
STDIN “connected” to keyboard by default,  
STDOUT connected to the console/terminal as  
well as STDERR



```
cssmuadm@lnx: /dev$ ls -l std*
lrwxrwxrwx 1 root root 15 May 1 18:55 stderr -> /proc/self/fd/2
lrwxrwxrwx 1 root root 15 May 1 18:55 stdin -> /proc/self/fd/0
lrwxrwxrwx 1 root root 15 May 1 18:55 stdout -> /proc/self/fd/1
cssmuadm@lnx: /dev$
```

# Standard Input and Output

**echo** command displays text using standard output

A terminal window with a title bar showing a file icon and the text 'cssmuadm@lnx: ~'. The terminal has a black background with white text. The first line shows the prompt 'cssmuadm@lnx:~\$' followed by the command 'echo "hello"'. The second line shows the output 'hello'. The third line shows the prompt 'cssmuadm@lnx:~\$' followed by a green cursor block.

```
cssmuadm@lnx:~$ echo "hello"
hello
cssmuadm@lnx:~$
```

# Standard Input and Output

Very simple **di** command:

- Reads from Standard Input
- Outputs it double times (\*:\*) to STDOUT
- Outputs message "STDERR OUTPUT: \*" to STDERR

```
cssmuadm@lnx: ~  
GNU nano 2.5.3 File: di.c  
#include <stdio.h>  
int main()  
{  
    char str[1024] = {0};  
    fscanf(stdin, "%s", (char*)&str);  
    fprintf(stdout, "%s:%s\n", str, str);  
    fprintf(stderr, "STDERR OUTPUT: '%s'\n", str);  
}
```

```
cssmuadm@lnx: ~  
cssmuadm@lnx:~$ di  
123  
123:123  
STDERR OUTPUT: '123'  
cssmuadm@lnx:~$
```

# Standard Input and Output: Pipes

Pipe (|) sends output of one program to input of another program. It's one of the most popular methods of interprocess communication in Linux

- Programs can work in parallel
- Pipes have limited size and allocated in *pipefs* (in memory file system)
- If pipe is full, write blocks until pipe has enough space to write, if pipe is empty, read blocks until there is some data to read from pipe

```
cssmuadm@lnx: ~  
cssmuadm@lnx:~$ echo "1" | di  
1:1  
STDERR OUTPUT: '1'  
cssmuadm@lnx:~$ echo "1" | di | di  
STDERR OUTPUT: '1'  
1:1:1:1  
STDERR OUTPUT: '1:1'  
cssmuadm@lnx:~$ echo "1" | di | di | di  
STDERR OUTPUT: '1'  
STDERR OUTPUT: '1:1'  
1:1:1:1:1:1:1:1  
STDERR OUTPUT: '1:1:1:1'  
cssmuadm@lnx:~$ echo "1" | di | di | di
```

```
cssmuadm@lnx: /proc/5274/fd  
-r--r--r-- 1 cssmuadm cssmuadm 0 May 9 16:09 status  
-r----- 1 cssmuadm cssmuadm 0 May 9 16:22 syscall  
dr-xr-xr-x 3 cssmuadm cssmuadm 0 May 9 16:22 task  
-r--r--r-- 1 cssmuadm cssmuadm 0 May 9 16:22 timers  
-rw-r--r-- 1 cssmuadm cssmuadm 0 May 9 16:22 uid_map  
-r--r--r-- 1 cssmuadm cssmuadm 0 May 9 16:22 wchan  
cssmuadm@lnx:/proc/5274$ cd fd  
cssmuadm@lnx:/proc/5274/fd$ ls -l  
total 0  
lrwx----- 1 cssmuadm cssmuadm 64 May 9 16:22 0 -> /dev/pts/0  
l-wx----- 1 cssmuadm cssmuadm 64 May 9 16:22 1 -> pipe:[201843]  
lrwx----- 1 cssmuadm cssmuadm 64 May 9 16:20 2 -> /dev/pts/0  
cssmuadm@lnx:/proc/5274/fd$
```

STDOUT  
redirected to pipe

# Standard Input and Output: Redirection

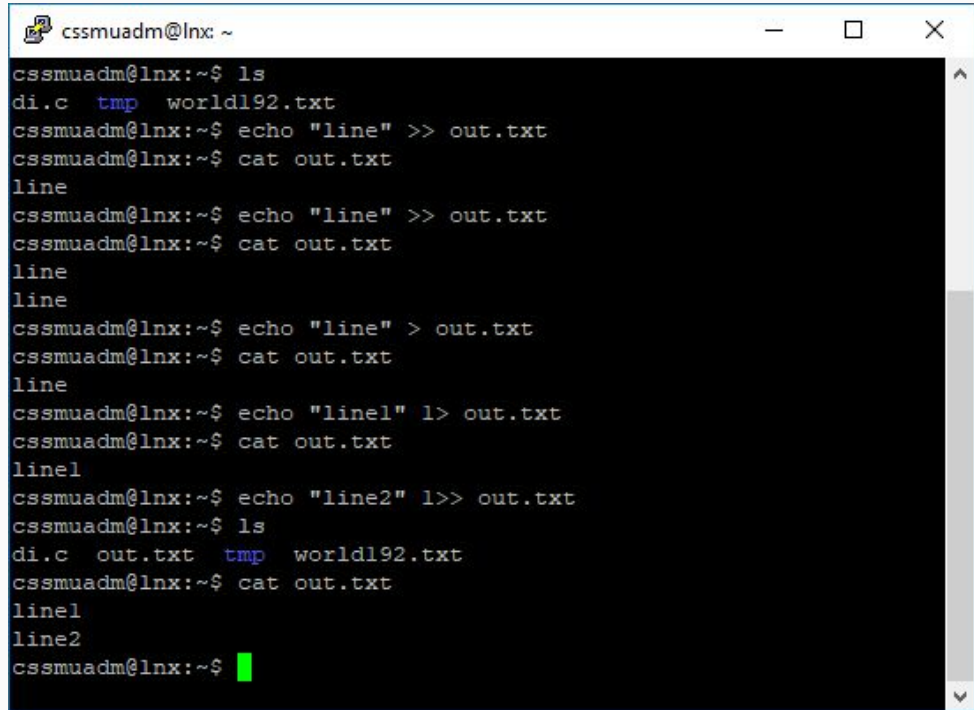
Redirect output (>) redirects output (STDOUT) to **file**. The same as **1>** as 1 stays for stdout file descriptor

```
cssmuadm@lnx: ~  
cssmuadm@lnx:~$ ls -l  
total 2444  
-rw-rw-r-- 1 cssmuadm cssmuadm 12 May 9 14:35 1  
-rwxrwxr-x 1 cssmuadm cssmuadm 8856 May 9 14:35 di  
-rw-rw-r-- 1 cssmuadm cssmuadm 187 May 9 14:35 di.c  
-rw-rw-r-- 1 cssmuadm cssmuadm 35 May 9 18:02 out.txt  
drwxrwxr-x 3 cssmuadm cssmuadm 4096 May 7 14:11 tmp  
-rw-r--r-- 1 cssmuadm cssmuadm 2473400 May 7 14:12 world192.txt  
cssmuadm@lnx:~$ ls -l > out.txt  
cssmuadm@lnx:~$ cat out.txt  
total 2440  
-rw-rw-r-- 1 cssmuadm cssmuadm 12 May 9 14:35 1  
-rwxrwxr-x 1 cssmuadm cssmuadm 8856 May 9 14:35 di  
-rw-rw-r-- 1 cssmuadm cssmuadm 187 May 9 14:35 di.c  
-rw-rw-r-- 1 cssmuadm cssmuadm 0 May 9 18:02 out.txt  
drwxrwxr-x 3 cssmuadm cssmuadm 4096 May 7 14:11 tmp  
-rw-r--r-- 1 cssmuadm cssmuadm 2473400 May 7 14:12 world192.txt  
cssmuadm@lnx:~$
```

# Standard Input and Output: Redirection

Double greater than sign (`>>`) also redirects STDOUT to a file. It creates a new file if not present, otherwise appends to it.

While single greater than sign (`>`) always **overwrites** file



```
cssmuadm@lnx: ~  
cssmuadm@lnx:~$ ls  
di.c tmp world192.txt  
cssmuadm@lnx:~$ echo "line" >> out.txt  
cssmuadm@lnx:~$ cat out.txt  
line  
cssmuadm@lnx:~$ echo "line" >> out.txt  
cssmuadm@lnx:~$ cat out.txt  
line  
line  
cssmuadm@lnx:~$ echo "line" > out.txt  
cssmuadm@lnx:~$ cat out.txt  
line  
cssmuadm@lnx:~$ echo "line1" 1> out.txt  
cssmuadm@lnx:~$ cat out.txt  
line1  
cssmuadm@lnx:~$ echo "line2" 1>> out.txt  
cssmuadm@lnx:~$ ls  
di.c out.txt tmp world192.txt  
cssmuadm@lnx:~$ cat out.txt  
line1  
line2  
cssmuadm@lnx:~$
```

# Standard Input and Output: Redirection

To redirect STDERR use **2>** (or **2>>** to append)

To redirect both STDOUT and STDERR use **&>**

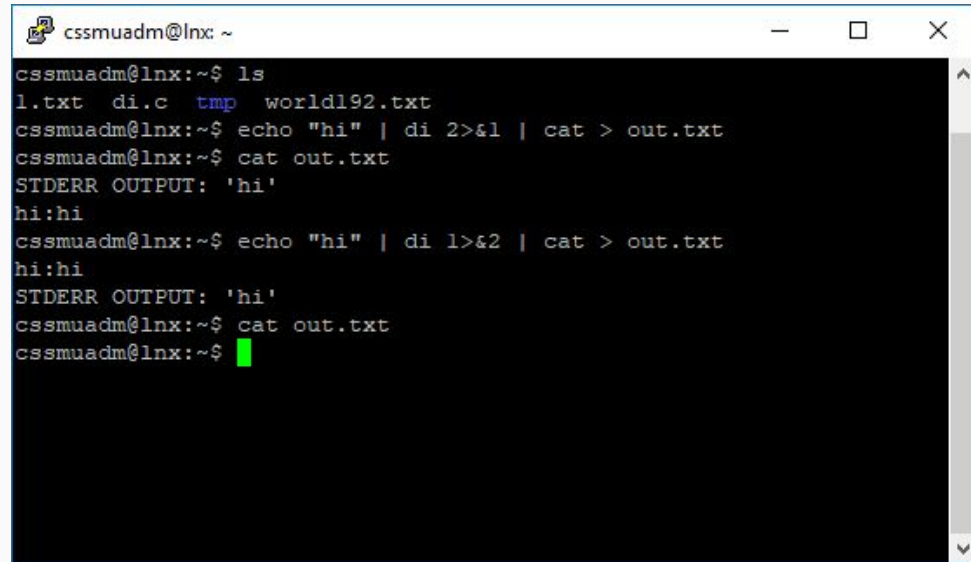
```
cssmuadm@lnx: ~  
cssmuadm@lnx:~$ echo "hi" | di  
hi:hi  
STDERR OUTPUT: 'hi'  
cssmuadm@lnx:~$ echo "hi" | di 1> out.txt  
STDERR OUTPUT: 'hi'  
cssmuadm@lnx:~$ cat out.txt  
hi:hi  
cssmuadm@lnx:~$ echo "hi" | di 2> out.txt  
hi:hi  
cssmuadm@lnx:~$ cat out.txt  
STDERR OUTPUT: 'hi'  
cssmuadm@lnx:~$ echo "hi" | di &> out.txt  
cssmuadm@lnx:~$ cat out.txt  
STDERR OUTPUT: 'hi'  
hi:hi  
cssmuadm@lnx:~$
```



# Standard Input and Output: Redirection

To redirect STDERR to STDOUT use  
**2>&1**

To redirect STDOUT to STDERR use  
**1>&2**

A terminal window titled 'cssmuadm@lnx: ~' with standard window controls. It shows a series of commands and their outputs. First, 'ls' lists files. Then, 'echo "hi" | di 2>&1 | cat > out.txt' is run, followed by 'cat out.txt' which outputs 'STDERR OUTPUT: 'hi'' and 'hi:hi'. Next, 'echo "hi" | di 1>&2 | cat > out.txt' is run, followed by 'cat out.txt' which outputs 'hi:hi' and 'STDERR OUTPUT: 'hi''. The prompt ends with a green cursor.

```
cssmuadm@lnx: ~  
cssmuadm@lnx:~$ ls  
l.txt  di.c  tmp  world192.txt  
cssmuadm@lnx:~$ echo "hi" | di 2>&1 | cat > out.txt  
cssmuadm@lnx:~$ cat out.txt  
STDERR OUTPUT: 'hi'  
hi:hi  
cssmuadm@lnx:~$ echo "hi" | di 1>&2 | cat > out.txt  
hi:hi  
STDERR OUTPUT: 'hi'  
cssmuadm@lnx:~$ cat out.txt  
cssmuadm@lnx:~$
```

# Standard Input and Output: Redirection

To redirect something to nowhere (useful in scripts), use redirection to **/dev/null**

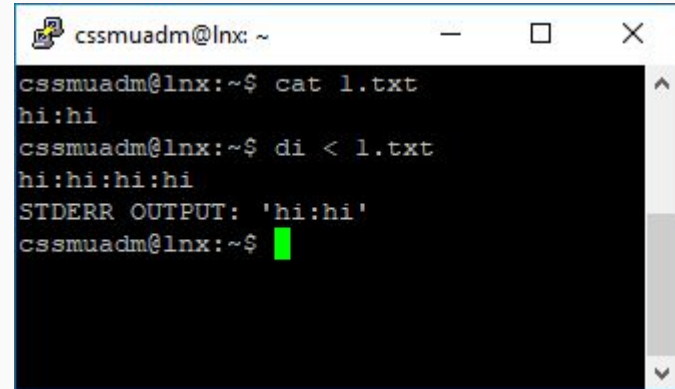
You can redirect STDERR to one file and STDOUT to another file

```
cssmuadm@lnx: ~  
cssmuadm@lnx:~$ ls  
di.c out.txt tmp world192.txt  
cssmuadm@lnx:~$ echo "hi" | di 2> /dev/null  
hi:hi  
cssmuadm@lnx:~$ echo "hi" | di 2> /dev/null 1> 1.txt  
cssmuadm@lnx:~$ cat 1.txt  
hi:hi  
cssmuadm@lnx:~$
```



# Standard Input and Output: Redirection

Less than sign (<) is used to accept input from file

A terminal window titled 'cssmuadm@lnx: ~' with standard window controls. It shows a sequence of commands and their outputs. First, 'cat 1.txt' is run, outputting 'hi:hi'. Then, 'di < 1.txt' is run, outputting 'hi:hi:hi:hi'. Finally, 'STDERR OUTPUT: 'hi:hi'' is displayed. The prompt 'cssmuadm@lnx:~\$' is followed by a green cursor.

```
cssmuadm@lnx:~$ cat 1.txt
hi:hi
cssmuadm@lnx:~$ di < 1.txt
hi:hi:hi:hi
STDERR OUTPUT: 'hi:hi'
cssmuadm@lnx:~$
```

# Standard Input and Output: Summary

- Pipes (`|`) are used to connect output of one program with input of another program
- Redirection is used to redirect STDOUT/STDERR to file (`1>`, `2>`) OR use contents of file as STDIN (`<`)

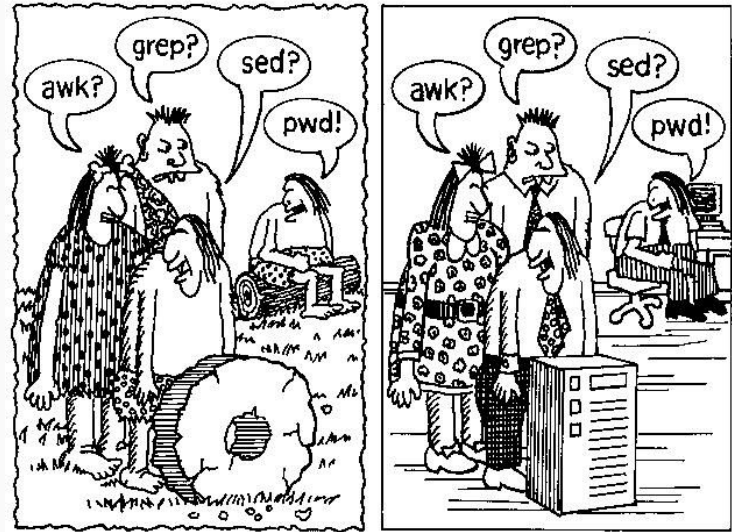
# Exercise

*List content of **/bin** directory to file with name **binlist.txt***

# Working with text, regular expressions

There is a set of programs that can get text data from STDIN (or as an argument) and transform this text data someway:

- sort
- grep
- awk
- tr
- sed
- shuf
- wc ...etc



# sort

**sort** - sort lines of text files

- **-n** does numerical sort rather than string comparison
- **-r** does reverse sort
- **-t** specifies field separator (e.g., *sort -t','*)
- **-k** sort by key (usually field number goes after key like *k2*)

```
root@lnx: /home/student
root@lnx:/home/student# ls | sort | head
a_cornish
a_hassan
a_sharma
a_sivaraman
b_ahuja
bb_verma
b_shree
ca_irfanoglu
c_mccavour
d_dsouzabhata
root@lnx:/home/student# ls | sort -r | head
z_yan
x_tong
v_govindan
v_boopathy
s_thapamagar
s_punnoli
s_padikar
s_magray
sk_channaveerabhadraiah
sh_sadi
root@lnx:/home/student#
```

# numeric sort vs lexicographical sort

```
cssmuadm@lnx: /bin
cssmuadm@lnx:/bin$ du -s * | sort -r | head
96    ps
96    gzip
96    df
84    dumpkeys
8     znew
8     zgrep
8     zdiff
8     lesspipe
8     gzexe
8     bzexe
cssmuadm@lnx:/bin$ du -s * | sort -r -n | head
1920   busybox
1016   bash
664    networkctl
648    systemctl
512    btrfs
488    journalctl
444    loginctl
440    udevadm
376    tar
368    ip
cssmuadm@lnx:/bin$
```



# Exercise

*Sort list of students (/home/students directory) by student **last name***

# grep

**grep** - print lines matching a pattern

Can work both as FILTER (accept STDIN) or  
with command line arguments

**-R** means recursively scan the directory from the  
command line argument

# grep

```
root@lnx: /home/student

root@lnx:/home/student# ls | grep an
a_hassan
a_sivaraman
ca_irfanoglu
dk_sambhwani
j_laplane
j_wang
kk_murugappan
p_bandyopadhyay
sk_channaveerabhadraiah
v_govindan
z_yan
root@lnx:/home/student# ls | grep an$
a_hassan
a_sivaraman
kk_murugappan
v_govindan
z_yan
root@lnx:/home/student# grep -R ERROR /var/log
/var/log/apport.log.1:ERROR: apport (pid 6549) Sat Jan 12 17:42:42 2019: called for pid 6548, signal 11, core limit 0,
dump mode 1
/var/log/apport.log.1:ERROR: apport (pid 6549) Sat Jan 12 17:42:42 2019: executable: /lib/klibc-k3La8MUnuzHQ0_kG8hokcG
AC0PA.so (command line "/lib/klibc-k3La8MUnuzHQ0_kG8hokcGAC0PA.so")
/var/log/apport.log.1:ERROR: apport (pid 6549) Sat Jan 12 17:42:42 2019: is_closing_session(): no DBUS_SESSION_BUS_ADD
RESS in environment
/var/log/apport.log.1:ERROR: apport (pid 6549) Sat Jan 12 17:42:42 2019: wrote report /var/crash/_lib_klibc-k3La8MUnuz
HQ0_kG8hokcGAC0PA.so.1028.crash
^C
root@lnx:/home/student#
```

# grep

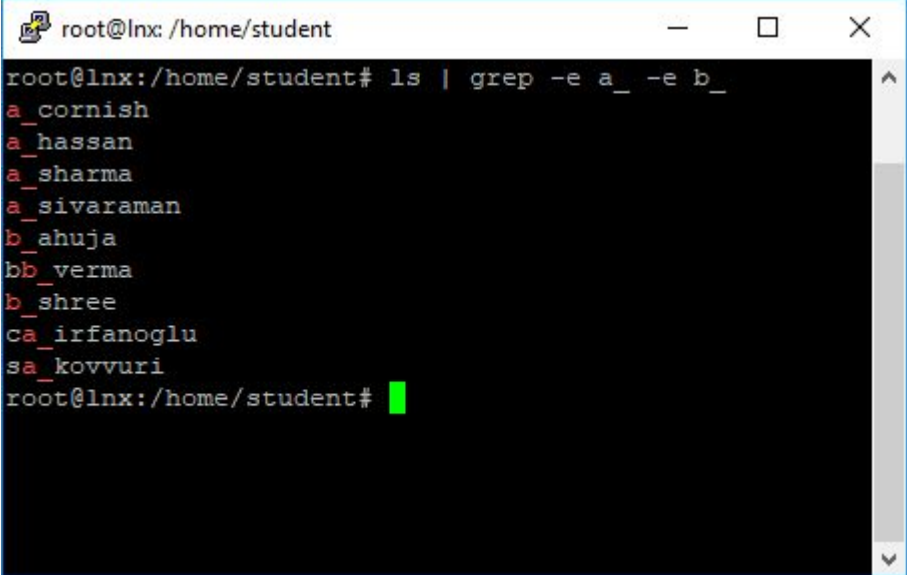
**-v** flag **INVERTS** match (finds lines NOT matching pattern)

```
root@lnx: /home/student
root@lnx:/home/student# ls | grep -v s_
a_cornish
a_hassan
a_sharma
a_sivaraman
b_ahuja
bb_verma
b_shree
ca_irfanoglu
c_mccavour
d_dsouzabhatt
dk_sambhwani
d_malone
g_singh
j_kour
j_laplane
j_wang
j_weston
kk_murugappan
m_fasuyi
mk_baria
mm_leong
mn_hussain
p_bandyopadhyay
q_cai
r_gupta
r_nomula
sa_kovvuri
sh_sadi
sk_channaveerabhadraiah
v_boopathy
v_govindan
x_tong
z_yan
root@lnx:/home/student#
```

# grep

**-e** PATTERN **-e** PATTERN **-e**  
PATTERN

Match many patterns at a time

A terminal window titled 'root@lnx: /home/student' with standard window controls. The command 'ls | grep -e a\_ -e b\_' is entered. The output lists files with names matching either 'a\_' or 'b\_': a\_cornish, a\_hassan, a\_sharma, a\_sivaraman, b\_ahuja, bb\_verma, b\_shree, ca\_irfanoglu, and sa\_kovvuri. The prompt 'root@lnx:/home/student#' is visible at the bottom with a green cursor.

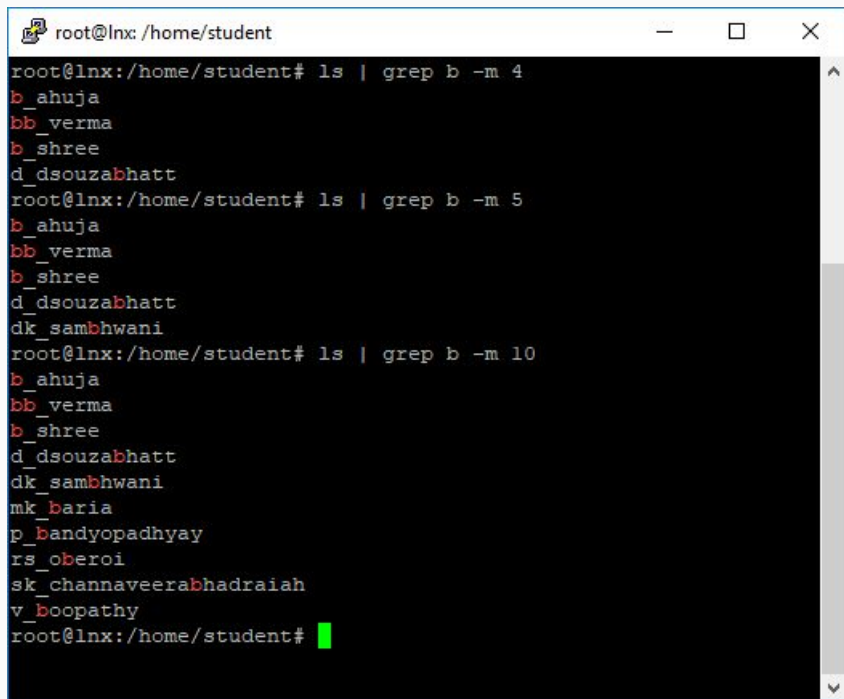
```
root@lnx: /home/student
root@lnx:/home/student# ls | grep -e a_ -e b_
a_cornish
a_hassan
a_sharma
a_sivaraman
b_ahuja
bb_verma
b_shree
ca_irfanoglu
sa_kovvuri
root@lnx:/home/student#
```

# grep as much as you need

```
root@lnx: /home/student
root@lnx:/home/student# ls | grep -e a_ -e b_
a_cornish
a_hassan
a_sharma
a_sivaraman
b_ahuja
bb_verma
b_shree
ca_irfanoglu
sa_kovvuri
root@lnx:/home/student# ls | grep -e a_ -e b_ | grep -v s
b_ahuja
bb_verma
ca_irfanoglu
root@lnx:/home/student#
```

# grep

**-m** flag LIMITS number of output lines per file

A terminal window titled 'root@lnx: /home/student' with standard window controls. It shows three sequential grep commands and their outputs. The first command 'ls | grep b -m 4' lists four names. The second 'ls | grep b -m 5' lists five names. The third 'ls | grep b -m 10' lists ten names. The output is color-coded: 'b' is red and the rest of the name is white.

```
root@lnx:/home/student# ls | grep b -m 4
b_ahuja
bb_verma
b_shree
d_dsouzabhatt
root@lnx:/home/student# ls | grep b -m 5
b_ahuja
bb_verma
b_shree
d_dsouzabhatt
dk_sambbhwani
root@lnx:/home/student# ls | grep b -m 10
b_ahuja
bb_verma
b_shree
d_dsouzabhatt
dk_sambbhwani
mk_baria
p_bandyopadhyay
rs_beroi
sk_channaveerabhadraiah
v_boopathy
root@lnx:/home/student#
```

# grep

-A N print N lines AFTER match

-B N print N lines BEFORE match

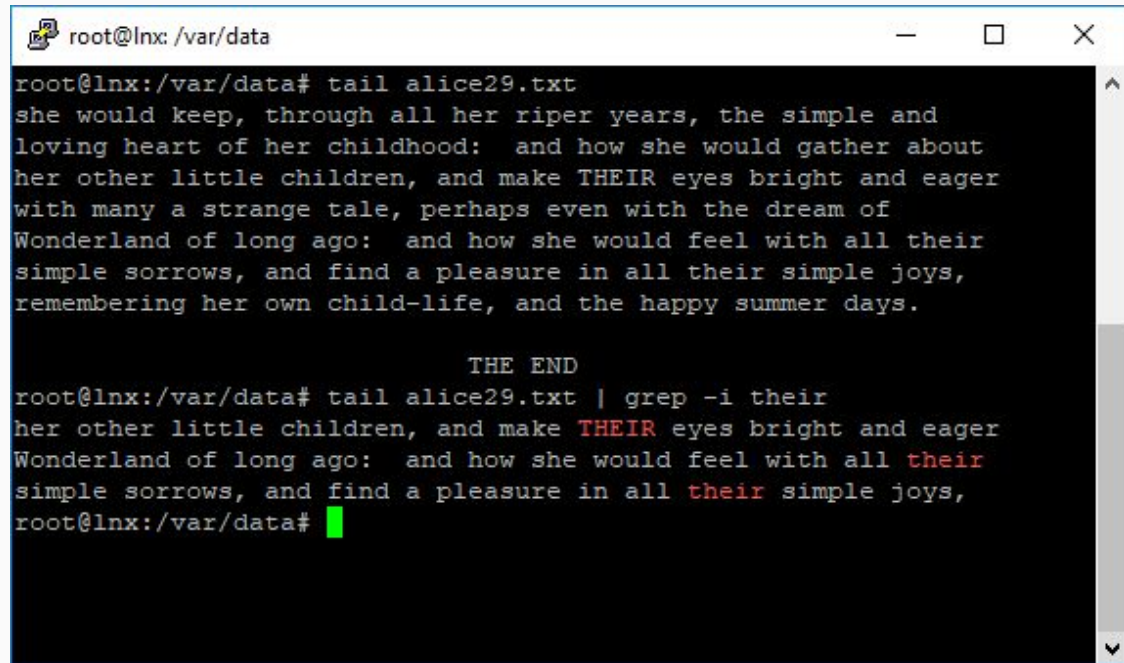
use case: analyzing logs

```
root@lnx: /home/student
root@lnx:/home/student# grep -A 2 -B 2 j_salvi /var/log/auth*
/var/log/auth.log-May 8 22:16:29 lnx systemd: pam_unix(systemd-user:session): session opened for user v_nagamanickam by (uid=0)
)
/var/log/auth.log-May 8 22:16:29 lnx su[29687]: pam_systemd(su:session): Cannot create session: Already running in a session
/var/log/auth.log-May 8 22:16:29 lnx sshd[29572]: Accepted password for j_salvi from 140.184.16.150 port 50464 ssh2
/var/log/auth.log-May 8 22:16:29 lnx sshd[29572]: pam_unix(sshd:session): session opened for user j_salvi by (uid=0)
/var/log/auth.log-May 8 22:16:29 lnx systemd: pam_unix(systemd-user:session): session opened for user j_salvi by (uid=0)
/var/log/auth.log-May 8 22:16:29 lnx systemd-logind[757]: New session 560 of user j_salvi.
/var/log/auth.log-May 8 22:16:41 lnx sshd[29802]: Invalid user Dauneo from 140.184.17.209
/var/log/auth.log-May 8 22:16:41 lnx sshd[29802]: input_userauth_request: invalid user Dauneo [preauth]
--
/var/log/auth.log-May 9 01:45:39 lnx systemd-logind[757]: Removed session 561.
/var/log/auth.log-May 9 01:45:39 lnx systemd: pam_unix(systemd-user:session): session closed for user r_mishra
/var/log/auth.log-May 9 01:48:23 lnx sshd[29572]: pam_unix(sshd:session): session closed for user j_salvi
/var/log/auth.log-May 9 01:48:23 lnx systemd-logind[757]: Removed session 560.
/var/log/auth.log-May 9 01:48:23 lnx systemd: pam_unix(systemd-user:session): session closed for user j_salvi
/var/log/auth.log-May 9 01:49:29 lnx sshd[29251]: pam_unix(sshd:session): session closed for user o_selere
/var/log/auth.log-May 9 01:49:29 lnx systemd-logind[757]: Removed session 551.
--
/var/log/auth.log-May 10 22:40:43 lnx sshd[21184]: pam_unix(sshd:session): session opened for user m_maruf by (uid=0)
/var/log/auth.log-May 10 22:40:43 lnx systemd-logind[757]: New session 779 of user m_maruf.
/var/log/auth.log-May 10 22:40:44 lnx sshd[21182]: Accepted password for j_salvi from 140.184.17.227 port 50874 ssh2
/var/log/auth.log-May 10 22:40:44 lnx sshd[21182]: pam_unix(sshd:session): session opened for user j_salvi by (uid=0)
/var/log/auth.log-May 10 22:40:44 lnx systemd-logind[757]: New session 780 of user j_salvi.
/var/log/auth.log-May 10 22:40:44 lnx systemd: pam_unix(systemd-user:session): session opened for user j_salvi by (uid=0)
/var/log/auth.log-May 10 22:40:54 lnx sshd[21258]: Accepted password for j_salvi from 140.184.17.227 port 50875 ssh2
/var/log/auth.log-May 10 22:40:54 lnx sshd[21258]: pam_unix(sshd:session): session opened for user j_salvi by (uid=0)
/var/log/auth.log-May 10 22:40:54 lnx systemd-logind[757]: New session 781 of user j_salvi.
/var/log/auth.log-May 10 22:41:27 lnx sshd[21386]: Accepted password for y_yin from 140.184.206.20 port 55180 ssh2
/var/log/auth.log-May 10 22:41:27 lnx sshd[21386]: pam_unix(sshd:session): session opened for user y_yin by (uid=0)
--
/var/log/auth.log-May 11 00:52:10 lnx systemd-logind[757]: Removed session 776.
/var/log/auth.log-May 11 00:55:03 lnx sshd[24046]: Connection closed by 104.131.202.23 port 42504 [preauth]
/var/log/auth.log-May 11 00:56:16 lnx sshd[21182]: pam_unix(sshd:session): session closed for user j_salvi
/var/log/auth.log-May 11 00:56:16 lnx systemd-logind[757]: Removed session 780.
/var/log/auth.log-May 11 00:58:01 lnx sshd[24050]: Did not receive identification string from 67.42.27.101
--
/var/log/auth.log-May 11 01:39:01 lnx CRON[24177]: pam_unix(cron:session): session opened for user root by (uid=0)
/var/log/auth.log-May 11 01:39:01 lnx CRON[24177]: pam_unix(cron:session): session closed for user root
/var/log/auth.log-May 11 01:42:08 lnx sshd[21258]: pam_unix(sshd:session): session closed for user j_salvi
/var/log/auth.log-May 11 01:42:08 lnx systemd-logind[757]: Removed session 781.
```



# case-insensitive grep

-i is case-insensitive

A terminal window titled 'root@lnx: /var/data' with standard window controls. It shows the output of 'tail alice29.txt' and a subsequent 'grep -i their' command. The grep results highlight 'THEIR' and 'their' in red text.

```
root@lnx:/var/data# tail alice29.txt
she would keep, through all her riper years, the simple and
loving heart of her childhood:  and how she would gather about
her other little children, and make THEIR eyes bright and eager
with many a strange tale, perhaps even with the dream of
Wonderland of long ago:  and how she would feel with all their
simple sorrows, and find a pleasure in all their simple joys,
remembering her own child-life, and the happy summer days.

                                THE END
root@lnx:/var/data# tail alice29.txt | grep -i their
her other little children, and make THEIR eyes bright and eager
Wonderland of long ago:  and how she would feel with all their
simple sorrows, and find a pleasure in all their simple joys,
root@lnx:/var/data#
```

# Exercise

*Using **grep**, find from which ip address you accessed the lnx server last time*

*logfiles: **/var/log/auth\*** (you have temporary read permissions)*

TIP: **sshd** keyword can help you to filter records that contain IP address info. Your **username** is also an obvious filter

# Regex basics

It's something like **"NORMAL TEXT + some special symbols"** that **MATCHES** text

STRING: BANANA

Regex 1: BANANA

Regex 1: ^BANANA\$

Regex 2: B.NA.A

Regex 3: BAN...

Regex 4: BO.\*A //does not match

Regex 5: B.\*A

Regex 6: .\*



*ANCIENT EGYPTIAN REGEXP*

geek and poke

[\[link\]](#)

# ^Regex\$

^ matches BEGINNING OF LINE

\$ matches END OF LINE

```
root@lnx: /home/student
root@lnx:/home/student# ls | grep a_
a_cornish
a_hassan
a_sharma
a_sivaraman
ca_irfanoglu
sa_kovvuri
root@lnx:/home/student# ls | grep ^a_
a_cornish
a_hassan
a_sharma
a_sivaraman
root@lnx:/home/student# ls | grep an
a_hassan
a_sivaraman
ca_irfanoglu
dk_sambhwani
j_laplante
j_wang
kk_murugappan
p_bandyopadhyay
sk_channaveerabhadraiah
v_govindan
z_yan
root@lnx:/home/student# ls | grep an$
a_hassan
a_sivaraman
kk_murugappan
v_govindan
z_yan
root@lnx:/home/student#
```

# \bRegex\b

\b is a word boundary

```
root@lnx: /var/data
```

```
root@lnx:/var/data# tail alice29.txt
she would keep, through all her riper years, the simple and
loving heart of her childhood: and how she would gather about
her other little children, and make THEIR eyes bright and eager
with many a strange tale, perhaps even with the dream of
Wonderland of long ago: and how she would feel with all their
simple sorrows, and find a pleasure in all their simple joys,
remembering her own child-life, and the happy summer days.
```

THE END

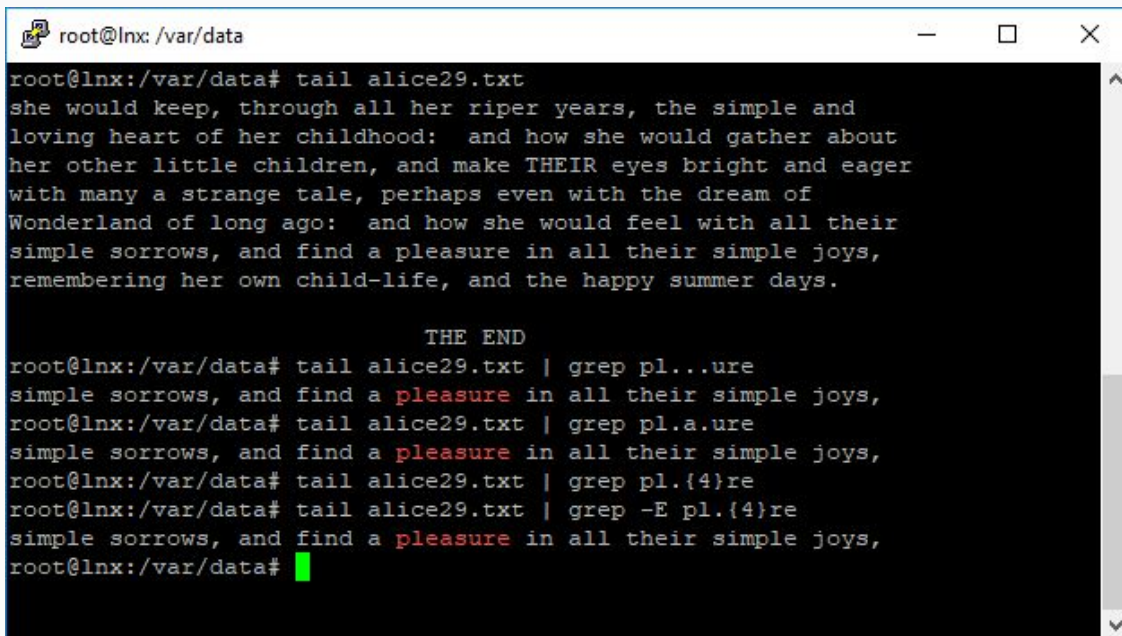
```
root@lnx:/var/data# tail alice29.txt | grep \bwould\b
she would keep, through all her riper years, the simple and
loving heart of her childhood: and how she would gather about
Wonderland of long ago: and how she would feel with all their
root@lnx:/var/data# tail alice29.txt | grep \bwould\b
she would keep, through all her riper years, the simple and
loving heart of her childhood: and how she would gather about
Wonderland of long ago: and how she would feel with all their
root@lnx:/var/data# tail alice29.txt | grep \bw\b
she would keep, through all her riper years, the simple and
loving heart of her childhood: and how she would gather about
Wonderland of long ago: and how she would feel with all their
root@lnx:/var/data#
```

# R..ex

. (DOT) matches ANY **single** character

You can use **N** dots to match any **N** characters

NOTE: -E flag stays for Extended, can always be used

A terminal window titled 'root@lnx: /var/data' with standard window controls. It shows a sequence of commands and their outputs. First, 'tail alice29.txt' displays a paragraph from a text file. Then, 'THE END' is printed. Finally, four 'grep' commands are used to search for the word 'pleasure' with different patterns: 'pl...ure', 'pl.a.ure', 'pl.{4}re', and '-E pl.{4}re'. The word 'pleasure' is highlighted in red in the output of each search.

```
root@lnx:/var/data# tail alice29.txt
she would keep, through all her riper years, the simple and
loving heart of her childhood: and how she would gather about
her other little children, and make THEIR eyes bright and eager
with many a strange tale, perhaps even with the dream of
Wonderland of long ago: and how she would feel with all their
simple sorrows, and find a pleasure in all their simple joys,
remembering her own child-life, and the happy summer days.

                                THE END
root@lnx:/var/data# tail alice29.txt | grep pl...ure
simple sorrows, and find a pleasure in all their simple joys,
root@lnx:/var/data# tail alice29.txt | grep pl.a.ure
simple sorrows, and find a pleasure in all their simple joys,
root@lnx:/var/data# tail alice29.txt | grep pl.{4}re
root@lnx:/var/data# tail alice29.txt | grep -E pl.{4}re
simple sorrows, and find a pleasure in all their simple joys,
root@lnx:/var/data#
```

# R+ege{1}xX\*

\* (STAR) works as a quantifier: matches ANY number of character **specified right before it**.

Examples:

- `.*` means any number of any character (matches everything)
- `a*` means any number of `a` (including zero)

+ (PLUS) also works as a quantifier: matches **one or more number** of character specified right before it

- `{N}` to match **exactly N** characters
- `{N,}` to match **N and more** characters
- `{,N}` to match **up to N** characters
- `{M,N}` to match **from M to N** characters

```
root@lnx: /home/student
root@lnx:/home/student# ls
a_cornish      dk_sambhwani  nm_leong       sk_channaveerabhadraiah
a_hassan       d_malone      mn_hussain     s_magray
a_sharma       g_singh       p_bandyopadhyay s_padikar
a_sivaraman    j_kour        q_cai          s_punnoli
b_ahuja        j_laplante    r_gupta        s_thapamagar
bb_verma       j_wang        r_nomula       v_boopathy
b_shree        j_weston      rs_oberoi      v_govindan
ca_irfanoglu   kk_murugappan sa_kovvuri     x_tong
c_mccavour     m_fasuyi      s_alachkar     z_yan
d_dsouzabhatt mk_baria      sh_sadi
root@lnx:/home/student# ls | grep ^s.*i$
sa_kovvuri
sh_sadi
s_punnoli
root@lnx:/home/student# ls | grep -E ^sk_chan*aveerabhadraiah$
sk_channaveerabhadraiah
root@lnx:/home/student# ls | grep -E ^sk_chan+aveerabhadraiah$
sk_channaveerabhadraiah
root@lnx:/home/student# ls | grep -E ^sk_channZ*aveerabhadraiah$
sk_channaveerabhadraiah
root@lnx:/home/student# ls | grep -E ^sk_channZ+aveerabhadraiah$
sk_channaveerabhadraiah
root@lnx:/home/student# ls | grep -E ^sk_chan{2}aveerabhadraiah$
sk_channaveerabhadraiah
root@lnx:/home/student#
```



# R[a-z]gex

[abc] set of symbols in SQUARE  
BRACKETS means that **any** of the  
symbols can be matched

Special cases:

[0-9] - any digit

[a-z] - any lowercase letter

[A-Z] -any uppercase letter

```
root@lnx: /home/student
root@lnx:/home/student# ls
a_cornish      dk_sambhwani  mm_leong       sk_channaveerabhadraiah
a_hassan       d_malone      mn_hussain     s_magray
a_sharma       g_singh       p_bandyopadhyay s_padikar
a_sivaraman    j_kour        q_cai          s_punnoli
b_ahuja        j_laplante    r_gupta        s_thapamagar
bb_verma       j_wang        r_nomula       v_boopathy
b_shree        j_weston      rs_oberoi      v_govindan
ca_irfanoglu   kk_murugappan sa_kovvuri     x_tong
c_mccavour     m_fasuyi      s_alachkar     z_yan
d_dsouzabhatt mk_baria      sh_sadi
root@lnx:/home/student# ls | grep -E ^[mq]+_.*$
m_fasuyi
mm_leong
q_cai
root@lnx:/home/student# ls | grep -E ^[mq]_.*$
m_fasuyi
q_cai
root@lnx:/home/student# ls | grep -E ^[xz]_.*$
x_tong
z_yan
root@lnx:/home/student#
```



# R(e)g\1x

() - everything between parentheses is a **capturing group**

*Q: Why you need them?*

*A: To reference what's between them again using a **backreference***

Backreference is a backslash followed by capturing group number (e.g., \1, \2, \3 ... \N) where N is number of capturing groups in your regex

```
root@lnx: /home/student
root@lnx:/home/student# ls
a_cornish      dk_sambhwani  mm_leong      sk_channaveerabhadraiah
a_hassan       d_malone     mn_hussain    s_magray
a_sharma       g_singh      p_bandyopadhyay s_padikar
a_sivaraman    j_kour       q_cai         s_punnoli
b_ahuja        j_laplante   r_gupta       s_thapamagar
bb_verma       j_wang       r_nomula      v_boopathy
b_shree        j_weston     rs_oberoi     v_govindan
ca_irfanoglu   kk_murugappan sa_kovvuri    x_tong
c_mccavour     m_fasuyi     s_alachkar    z_yan
d_dsouzabhatt mk_baria     sh_sadi
root@lnx:/home/student# ls | grep -E 's_th(.)p\lm\lg\lr'
s_thapamagar
root@lnx:/home/student#
```

# Regex\\.

What if we need to match a “**special**” character that have meaning in regex (e.g., “.”, “\*”, “\$”, “?”, “^”)?

Just escape them with \\

```
cssmuadm@lnx: ~  
cssmuadm@lnx:~$ cat test.doc  
hello hello . how ORE you..?  
How much $ do you have?  
hehe ^ ^  
cssmuadm@lnx:~$ cat test.doc | grep -E .  
hello hello . how ORE you..?  
How much $ do you have?  
hehe ^ ^  
cssmuadm@lnx:~$ cat test.doc | grep -E \\.  
hello hello . how ORE you..?  
cssmuadm@lnx:~$ cat test.doc | grep -E $  
hello hello . how ORE you..?  
How much $ do you have?  
hehe ^ ^  
cssmuadm@lnx:~$ cat test.doc | grep -E \\$  
How much $ do you have?  
cssmuadm@lnx:~$ cat test.doc | grep -E \\^  
hehe ^ ^  
cssmuadm@lnx:~$
```

# Exercise

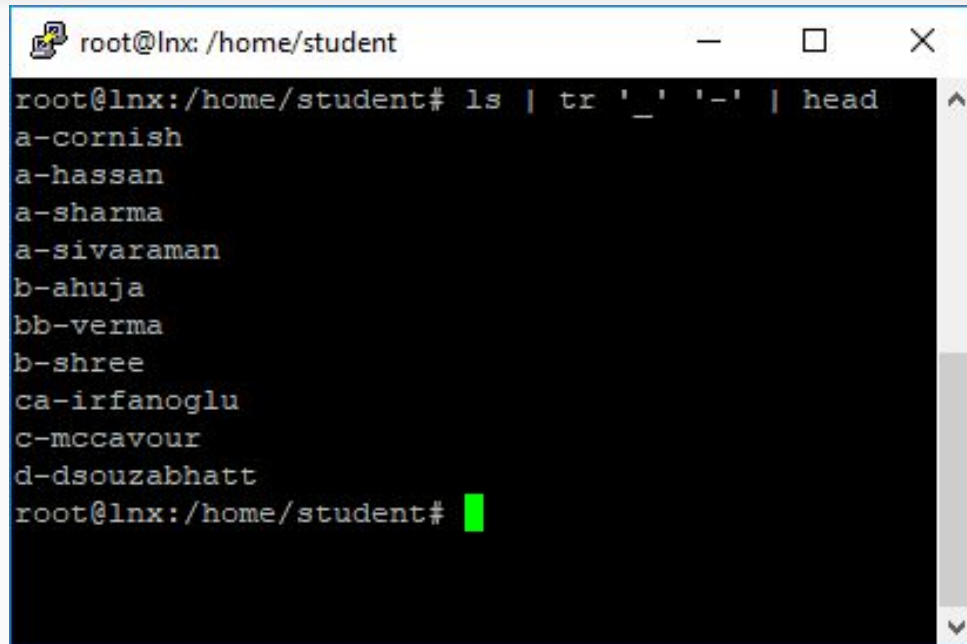
*Go to /home/student*

- *Use **ls + grep** output your and your neighbour's account together*
- *Using **ls + grep** find all students with **double letters** in their last name*

# tr

**tr** - translate or delete characters

tr [character to replace] [replacement]



```
root@lnx: /home/student
root@lnx:/home/student# ls | tr '_' '-' | head
a-cornish
a-hassan
a-sharma
a-sivaraman
b-ahuja
bb-verma
b-shree
ca-irfanoglu
c-mccavour
d-dsouzabhatt
root@lnx:/home/student#
```

# tr

```
root@lnx: /home/student
root@lnx:/home/student# ls | tr '_' '-' | head
a-cornish
a-hassan
a-sharma
a-sivaraman
b-ahuja
bb-verma
b-shree
ca-irfanoglu
c-mccavour
d-dsouzabhatter
root@lnx:/home/student# ls | tr '_' '-' | head | tr '[abc]' '_'
_ _ornish
_ h _ss _n
_ sh _rm
_ _siv _r _m _n
_ _ huj _
_ _verm _
_ _shree
_ _irf _noglu
_ _m _ _vour
d-dsouz _ _h _tt
root@lnx:/home/student#
```

# tr

tr -d [abcd]

-d means DELETE

```
root@lnx: /home/student
root@lnx:/home/student# ls | tr -d '_' | tail
shsadi
skchannaveerabhadraiah
smagray
spadikar
spunnoli
sthapamagar
vboopathy
vgovindan
xtong
zyan
root@lnx:/home/student#
```

# Exercise

*Go to /home/student*

*Use **ls + grep** to get your account and **tr** to transform your account name to UPPERCASE*

*Hint: [:lower:] is all lowercase and [:upper:] is all uppercase*

# 's/sad/sed/g'

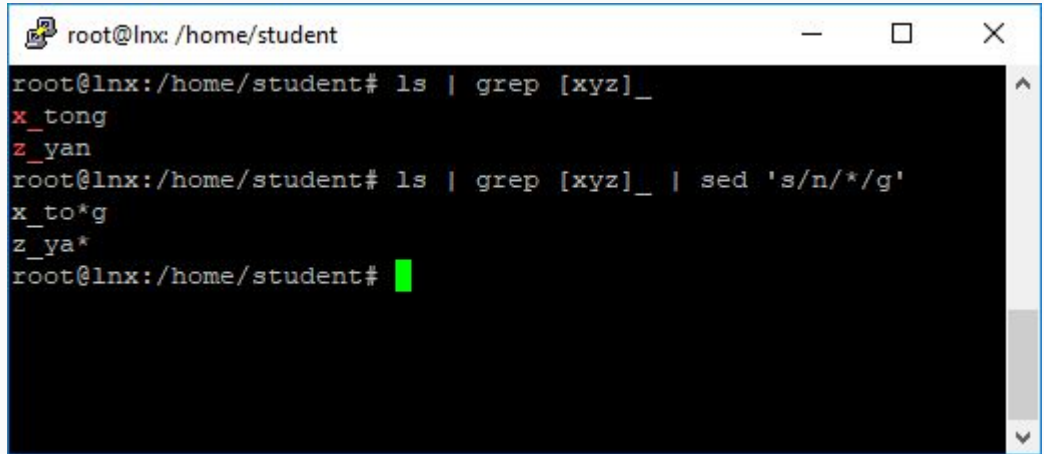
**sed** - stream editor for filtering and transforming

Common structure:

**sed -r 's/to\_replace/replacement/g'**

**-r** tells to use extended regular expressions

**g** at the end tells to replace **all** occurrences rather than just first one

A terminal window titled 'root@lnx: /home/student' with standard window controls. It shows a sequence of commands and their outputs. First, 'ls | grep [xyz]\_' lists files 'x\_tong' and 'z\_yan'. Then, the same command is followed by '| sed 's/n/\*/g'', which transforms the output to 'x\_to\*g' and 'z\_ya\*'.

```
root@lnx:/home/student# ls | grep [xyz]_
x_tong
z_yan
root@lnx:/home/student# ls | grep [xyz]_ | sed 's/n/*/g'
x_to*g
z_ya*
root@lnx:/home/student#
```



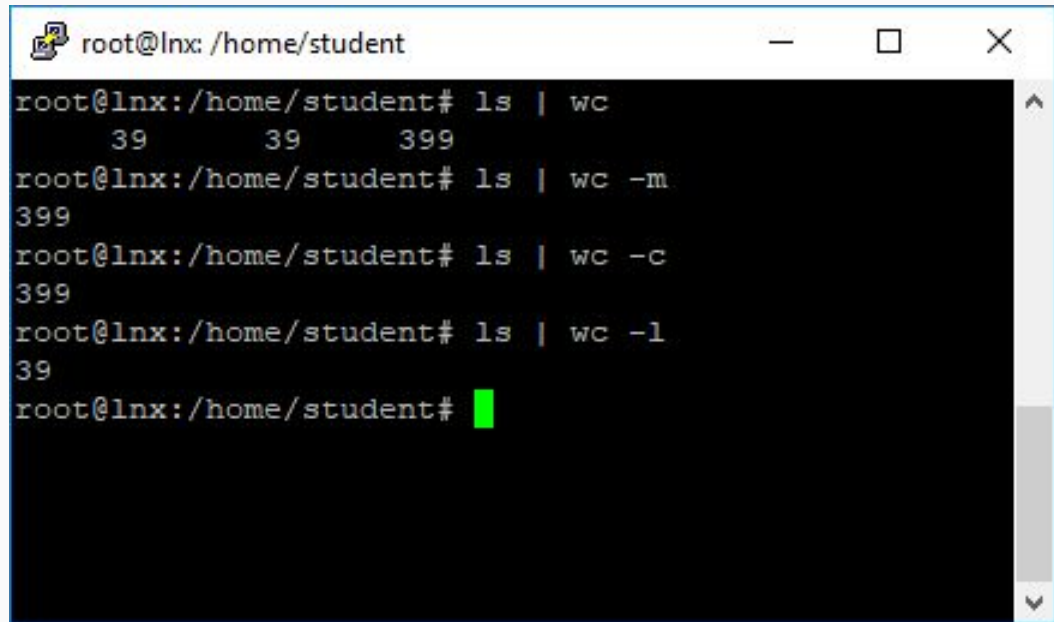
's/sa(.)/se\1/g'

```
root@lnx: /home/student
root@lnx:/home/student# ls | tail
sh_sadi
sk_channaveerabhadraiah
s_magray
s_padikar
s_punnoli
s_thapamagar
v_boopathy
v_govindan
x_tong
z_yan
root@lnx:/home/student# ls | tail | sed -r 's/^(.*)_(.*)$/\2_\1/g'
sadi_sh
channaveerabhadraiah_sk
magray_s
padikar_s
punnoli_s
thapamagar_s
boopathy_v
govindan_v
tong_x
yan_z
root@lnx:/home/student#
```

# wc != restroom

**wc** - print newline, word, and byte counts for each file

- c** --bytes      print the byte counts
- m** --chars     print the character counts
- l** --lines      print the newline counts

A terminal window titled 'root@lnx: /home/student' with standard window controls. It shows a series of commands and their outputs. The first command 'ls | wc' outputs '39 39 399'. The second 'ls | wc -m' outputs '399'. The third 'ls | wc -c' outputs '399'. The fourth 'ls | wc -l' outputs '39'. The prompt is then followed by a green cursor.

```
root@lnx: /home/student# ls | wc
 39    39   399
root@lnx: /home/student# ls | wc -m
399
root@lnx: /home/student# ls | wc -c
399
root@lnx: /home/student# ls | wc -l
39
root@lnx: /home/student#
```

# shuf

**shuf** - generate random permutations (shuffle lines)

```
root@lnx: /home/student
root@lnx:/home/student# ls | shuf | head
sa_kovvuri
j_wang
c_mccavour
d_dsouzabhatt
p_bandyopadhyay
ca_irfanoglu
z_yan
d_malone
s_punnoli
b_ahuja
root@lnx:/home/student# ls | shuf | head
sa_kovvuri
ca_irfanoglu
b_shree
x_tong
s_punnoli
v_boopathy
c_mccavour
q_cai
g_singh
sh_sadi
root@lnx:/home/student#
```

# split

**split** - split a file into pieces

**-l** flag: number of lines in splitted files

**-n** flag: number of chunks. Useful with option **-n l/N** where *N* number of files: will split into *N* files without splitting lines/records

```
root@lnx: ~/split
root@lnx:~# mkdir split
root@lnx:~# cd split/
root@lnx:~/split# ls /home/student | split -l 10
root@lnx:~/split# ls -l
total 16
-rw-r--r-- 1 root root 103 Jan 17 17:34 xaa
-rw-r--r-- 1 root root 96 Jan 17 17:34 xab
-rw-r--r-- 1 root root 99 Jan 17 17:34 xac
-rw-r--r-- 1 root root 101 Jan 17 17:34 xad
root@lnx:~/split# cat xaa
a_cornish
a_hassan
a_sharma
a_sivaraman
b_ahuja
bb_verma
b_shree
ca_irfanoglu
c_mccavour
d_dsouzabhath
root@lnx:~/split# wc -l *
10 xaa
10 xab
10 xac
9 xad
39 total
root@lnx:~/split#
```

# awk

"**AWK** is a language for processing text files. A file is treated as a sequence of records, and by default each line is a record. Each line is broken up into a sequence of fields, so we can think of the first word in a line as the first field, the second word as the second field, and so on. **An AWK program is a sequence of pattern-action statements.** AWK reads the input a line at a time. A line is scanned for each pattern in the program, and for each pattern that matches, the associated action is executed." - Alfred V. Aho<sup>[1]</sup>

Most used feature:

**awk '{print \$1, \$2}'**

where \$1 \$2 are numbers of columns/fields

Flag -F',' is used to specify field separator

```
root@lnx: /home/student
root@lnx:/home/student# ls -l | tail -5
drwxr-x--x 3 s_thapamagar      s_thapamagar      4096 Jan 14 21:20 s_thapamagar
drwxr-x--x 2 v_boopathy        v_boopathy        4096 Jan 14 13:18 v_boopathy
drwxr-x--x 2 v_govindan        v_govindan        4096 Jan 14 13:18 v_govindan
drwxr-x--x 2 x_tong            x_tong            4096 Jan 17 14:11 x_tong
drwxr-x--x 2 z_yan             z_yan             4096 Jan 14 13:18 z_yan
root@lnx:/home/student# ls -l | tail -5 | awk '{print $3}'
s_thapamagar
v_boopathy
v_govindan
x_tong
z_yan
root@lnx:/home/student# ls -l | tail -5 | awk '{print $4}'
s_thapamagar
v_boopathy
v_govindan
x_tong
z_yan
root@lnx:/home/student# ls -l | tail -5 | awk '{print $5}'
4096
4096
4096
4096
4096
root@lnx:/home/student#
```

exit